

BASELINE AI WORKFORCE OS

14-Day Pilot SOP

Mission Control supervises · Baseline OS coordinates · Hermes / OpenClaw / Claude Code execute

Document: onboarding-14-day-pilot-sop.pdf

Source: docs/onboarding/14_DAY_PILOT_SOP.md

Generated: 2026-05-29

Every paid pilot follows this exact 14-day arc. The goal is one measurable business outcome by day 14, not a tour of every panel. If we can't show ROI by day 14, we don't ask for the upgrade.

Pilot fundamentals

Item	Value
Length	14 calendar days
Price	\$1 — credit-card-on-file (sets up the upgrade path)
Headcount	1 operator, up to 3 AI employees
Verticals	one of the 9 launch verticals
Success metric	one operator-chosen primary KPI (hours saved OR \$ recovered OR conversion lift)
Exit	upgrade to paid plan, downgrade to free demo, or off-board cleanly

Day-by-day cadence

Day	Owner	Step	Touchpoint	Output
0	Sales	Kickoff call (45 min)	live video	signed pilot agreement, primary KPI chosen, vertical chosen
1	Customer Success	Data collection	async checklist	tools list, SOPs uploaded, sample data shared
2	CS	Memory setup	live (30 min)	Obsidian / Notion connected, top 10 docs indexed
3	CS	Workflow mapping	live (60 min)	3 candidate workflows scoped, picked one for week-1 mission
4	CS	Runtime onboarding	live (45 min)	Hermes / OpenClaw / Claude Code installed; harness passes
5	Operator	First mission launched	async	Swarm Mode run completes; verification judges green
6	Operator	Operator works alongside AI	async	first daily standup digest sent
7	CS	Week-1 review (30 min)	live	3 wins / 1 blocker / 1 ask documented
8–10	Operator	Productive use	async	AI workforce runs every business day
11	CS	ROI capture call (30 min)	live	hours-saved + \$-impact baseline measured; case study draft started
12	CS	Approval & escalation review	live (30 min)	every approval-gated action logged & reviewed
13	Sales + CS	ROI review + upgrade conversation	live (45 min)	conversion decision, plan tier chosen, invoice issued OR off-board scheduled
14	Sales	Pilot close	async	converted to paid OR off-boarded with clean shutdown

Day 0 — Kickoff (45 minutes)

Pre-call: customer signed the pilot agreement and provided billing details for the \$1 setup.

Agenda:

1. (5 min) Welcome + remind them what a 14-day pilot looks like.
2. (10 min) Pick the primary KPI. Only one. Choose from:
 - Hours saved per week by a specific person/role
 - \$ recovered from a specific revenue leak (e.g. aging AR)
 - Conversion lift on a specific funnel (e.g. lead → booked)
 - Cycle time on a specific workflow (e.g. estimate → sent)
3. (10 min) Pick the vertical template (PM / GC / Home Services / Real Estate / Mortgage / CPA / Law / Marketing Agency / AI Agency).
4. (10 min) Walk through the guided demo for that vertical so they see the AI workforce they're about to deploy.
5. (10 min) Schedule the next four sessions (days 2, 4, 7, 11, 13).
Send calendar invites before the call ends.

Artifacts produced:

- Pilot ticket in `pilots/<customer-id>.md` (see template at the bottom of this doc)
 - Calendar invites sent
 - Slack / WhatsApp channel created (`#pilot-<customer-slug>`)
-

Day 1 — Data collection (async)

Send the customer this single checklist. Don't move on until every box is checked.

- ☐ Company name, primary domain, billing email
- ☐ Tools list: CRM, calendar, voice, payments, docs/wiki
- ☐ 3 SOPs uploaded (PDF or doc link)
- ☐ Sample dataset (50 rows minimum) for the primary workflow
- ☐ Compliance constraints (HIPAA / GLBA / PCI / state regs / TCPA / CAN-SPAM)
- ☐ Approval policy: which actions require human sign-off
- ☐ Hours/week the operator can invest in week 1

This is gated. No data → no memory setup. No memory setup → no mission.

Day 2 — Memory setup (30 minutes live)

1. Connect the source — Notion or Obsidian (both work).
2. Index the top 10 docs the customer says drive their decisions.
3. Verify a query: ask Mission Control to recall something only the docs would know. Confirm citation shows in the response.

4. Document any documents the customer refuses to share (legal, HR, PII). Note them as out-of-scope.

Acceptance: a memory citation renders inside an AI employee card on the customer's workspace.

Day 3 — Workflow mapping (60 minutes live)

Whiteboard exercise:

1. List every recurring task the operator touches each day.
2. Mark each one with:
 - A — should be automated
 - S — should be supervised by AI but human-decides
 - H — human only
3. From the A list, pick 3 candidate workflows scored by:
 - frequency × time-saved × reversibility
4. Pick one for the week-1 mission. The other two land in the week-2 backlog.

Acceptance: one workflow named, scoped to ≤ 5 steps, owned by ≤ 2 AI employees.

Day 4 — Runtime onboarding (45 minutes live)

1. Install whichever runtime is appropriate:
 - Hermes — strategy + memory recall
 - OpenClaw / OpenCode — execution
 - Claude Code — implementation
2. Run `./scripts/runtime-validate.sh --runtime <id>` against the customer's Mission Control instance.
3. All 5 stages must PASS. If they don't, fix the runtime config (not Mission Control) and re-run.
4. Confirm the Runtime Validation panel shows a green band.

Acceptance: runtime validation harness prints `PASS` for every stage.

Day 5 — First mission (async)

Operator runs Swarm Mode in `/app/workflows/swarm` with a mission matching the week-1 workflow. Verification judges must return PASS before the run is allowed to persist.

If judges return `attention`:

- review the flagged item with the operator
- approve or revise
- re-run

Acceptance: one persisted run with judges all green.

Day 6 — Operator works alongside AI (async)

Operator uses Mission Control daily. Daily standup digest goes to their inbox at 7am local. CS monitors via `/app/runtime-validation` and the activity feed.

CS pings the operator if:

- no heartbeat in 24 hours
 - approval queue grows past 5 pending items
 - any agent shows `needs-attention` for > 2 hours
-

Day 7 — Week-1 review (30 minutes live)

Run through:

1. 3 wins the AI workforce delivered this week (with \$ or hours).
2. 1 blocker that stopped them from getting a 4th win.
3. 1 ask — what feature, integration, or skill they need next.

Update the pilot ticket. Adjust the week-2 backlog if needed.

Day 8–10 — Productive use (async)

CS touches the channel every 24 hours, lightweight. Operator drives.

Day 11 — ROI capture (30 minutes live)

Numbers, not vibes:

Metric	Baseline (pre-pilot)	Pilot result	Delta
Hours / week on the workflow	X	Y	Z saved
\$ recovered (AR / no-shows / leakage)	X	Y	+Z
Conversion (leads → booked, etc.)	X%	Y%	+Z pts
Cycle time	X days	Y days	-Z days

Start the case study draft (1 paragraph + 3 bullets).

Day 12 — Approval & escalation review (30 minutes live)

Open `/app/approvals` and walk the operator through every gated action since day 5. Confirm:

- they understood why each one paused
- the approval decision was correct in hindsight
- no action ran without their approval that should have paused

This is the trust transfer. The pilot doesn't convert unless the operator trusts the supervision layer.

Day 13 — ROI review + upgrade (45 minutes live)

Open the ROI capture from day 11 + the wins from day 7. Present the upgrade options:

Tier	Monthly	Best for
Operator	\$99	1 user, up to 5 AI employees, 1 runtime
Team	\$299	up to 5 users, unlimited AI employees, 2 runtimes
Workspace	\$799	up to 25 users, all runtimes, priority support

Ask the close question: "Based on the ROI you captured, which tier matches what you need next month?"

Outcomes:

- Convert → upgrade — invoice issued, billing flipped from \$1 pilot to paid plan, kickoff next week's expansion.
 - Convert → wait — keep at pilot tier for \$99/mo until they expand. Re-engage in 14 days.
 - No convert — schedule day-14 clean shutdown. Get a one-paragraph testimonial if possible.
-

Day 14 — Pilot close (async)

If converted:

- email confirmation of plan tier + first month's invoice
- schedule expansion call (next workflow live by day +14)
- enable the customer in `/pilots` as "converted"

If off-boarded:

- Disable the workspace (read-only for 30 days, then purge)
 - Export their memory + workflow definitions and email to them
 - Send a 4-question exit survey
-

Pilot ticket template — copy to `pilots/<customer-id>.md`

```
# Pilot - <Customer Name>
Started: YYYY-MM-DD
Vertical: <pm|gc|home-services|real-estate|mortgage|cpa|law-firm|marketing-agency|ai-agency>
Primary KPI: <hours-saved | dollars-recovered | conversion-lift | cycle-time>
Owner: <CS name>
Sales: <Sales name>

## Day 0 - Kickoff
- KPI: ...
- Workflow shortlist: 1) ... 2) ... 3) ...

## Day 1 - Data
- Tools: ...
- SOPs uploaded: ...
- Compliance: ...

## Day 2 - Memory
- Source: Notion | Obsidian
- Indexed docs: <count>
- First citation: <link>

## Day 3 - Workflow mapping
- Week-1 mission: ...
- Week-2 backlog: ...

## Day 4 - Runtime
- Runtime: hermes | openclaw | claude
- Harness output: PASS / FAIL (paste log)

## Day 5 - First mission
- Mission link: /app/workflows/swarm/<run-id>
- Judge verdicts: pass / attention / fail
- Deliverables: ...

## Day 7 - Week-1 review
- Wins: 1) ... 2) ... 3) ...
- Blocker: ...
- Ask: ...

## Day 11 - ROI capture
- Hours saved/week: ...
- $ recovered: ...
- Conversion lift: ...

## Day 13 - Upgrade conversation
- Tier offered: ...
- Decision: convert | wait | no-convert
- Notes: ...

## Day 14 - Close
```

- Outcome: ...

- Next milestone: ...

Pilot health flags (early warning)

CS reviews these every morning across the active pilot cohort:

Flag	Trigger	Action
 Quiet	no operator activity > 24h	Slack ping the operator
 Memory drift	citations missing on 2+ employee cards	Re-index a doc on a live call
 Stalled	week-1 mission not launched by day 6	Sales joins the next call
 Lost trust	operator rejects 2+ approvals in a row	Pause AI workforce, schedule trust-rebuild call
 Runtime down	heartbeat absent > 4h	Runtime install call, re-run harness

A pilot ends week-1 with two or more  flags → CS escalates to sales for a re-scope. Don't try to convert a broken pilot.

Conversion benchmarks (internal)

These are targets, not promises:

- 80% of pilots complete the 14 days (don't abandon).
- 60% of completed pilots convert to a paid plan.
- 70% of conversions choose Team (\$299) on first upgrade.
- < 14 days median from kickoff to first measurable win.

If we drop below 50% conversion for two consecutive cohorts, the SOP gets a hard review. The number to watch is time to first measurable win. Below 7 days → conversion stays high. Above 10 days → conversion collapses.