

BASELINE AI WORKFORCE OS

AI Agency — Sales Playbook

Mission Control supervises · Baseline OS coordinates · Hermes / OpenClaw / Claude Code execute

Document: `playbook-ai-agency.pdf`

Source: `docs/sales/ai-agency.md`

Generated: 2026-05-29

AI Workforce OS for AI / automation agencies. Deploy an AI workforce per client in a day. Supervise dozens of client workspaces from one Mission Control.

1. One-page outline

Headline: AI Workforce OS for AI Agencies

Sub: Stop hand-rolling each client. Deploy production AI workforces in a day. Supervise every client account from one command center. Scale.

The 3 wins:

- New client AI workforce live in 1 day (not 4 weeks)
- One dashboard for every active client engagement
- Per-client memory, per-client billing, per-client approvals

Star employee: Mission Control Supervisor — the agency's delivery lead, but as an AI employee. Orchestrates per-client workspaces, raises blockers, prepares the weekly client report.

The roster (4 AI employees, agency-side):

Role	What it does	Memory source
Mission Control Supervisor	Multi-client orchestration	Agency playbook + client roster
Sales Follow-Up	Pipeline → onboarding handoff	CRM + proposal templates
Customer Intake	New client kickoff + data collection	Agency intake SOP
Reporting Bot	Weekly client outcome reports	Per-client metrics

The CTA:

1. Live demo: mission.baselineautomations.com/?demo=ai-agency
2. \$1, 14-day pilot (agency operator, 1 test client)
3. Convert at \$799/mo Workspace

2. Sales sheet

AI agencies have a delivery problem masquerading as a sales problem. Every client wants their own AI workforce. Every client gets a

hand-rolled stack. Five clients in, you're a delivery team, not an agency.

Baseline gives agencies the operating system:

- Per-client workspaces with isolated memory, billing, and approvals. One client never sees another.
- Pre-built vertical templates (PM, GC, Home Services, Real Estate, Mortgage, CPA, Law Firm, Marketing Agency) drop in. Customize the SOP, not the stack.
- Mission Control Supervisor orchestrates across clients. Weekly client reports drafted automatically. Per-client P&L.
- Demo share links mint instantly. A salesperson on a call hands over a watermarked, signed demo with the prospect's name in 10 seconds.
- Runtime supervision for Hermes, OpenClaw, Claude Code. Install once at the agency level, redeploy per client.

Mission Control supervises. Baseline OS coordinates. Hermes / OpenClaw / Claude Code execute. Your agency runs production AI operations, not custom builds.

Trial: \$1 / 14 days with one test client. Tier: \$799/mo
Workspace for up to 25 client workspaces.

3. Discovery script

1. How many active client engagements is the agency running?
 2. What's your average client onboarding time today?
 3. Are you hand-rolling client stacks or using a platform?
 4. Per-client billing / approvals / memory — separated or shared?
 5. What's your average client MRR?
 6. How are sales handoffs going from your sales rep to delivery?
 7. Weekly client reporting — manual or automated?
 8. Last client churn — root cause?
 9. Are you supervising AI runtimes already (Hermes / OpenClaw / Claude Code)?
 10. \$1, 14-day pilot on one test client, full agency workspace?
-

4. Objection handling

Objection	Response	Proof
"We built our own platform."	"Great. Use it for differentiated work. Use ours for the 80% that's the same across clients. We do per-client memory, billing, approvals, and runtime supervision — focus your engineers on the 20% that wins clients."	Demo workspace isolation
"Clients want bespoke."	"Vertical templates are starting points, not endings. Configure the SOP per client. Memory is per-client. Outputs are bespoke."	Show vertical template customization
"\$799/mo is more than a contractor."	"Per-client cost on Workspace tier drops to ~\$32/mo if you run 25 clients. Cheaper than a contractor's lunch budget."	Pricing math
"We use CrewAI / LangGraph / AutoGen."	"We're not a generic agent framework — we're the supervision + orchestration layer above the execution runtimes. CrewAI is welcome to run inside an OpenClaw-equivalent slot."	Architecture diagram
"Compliance — client data isolation."	"Per-client workspace ID, per-client memory namespace, per-client billing ledger. Audited per touch. SOC 2 path in progress."	Security doc
"We resell to clients — branding?"	"White-label is on the Workspace tier roadmap. For pilot, run on Baseline brand; switch on white-label when it ships."	Roadmap

5. Follow-up sequence

Day	Channel	Asset
0	Email	Recap + signed demo + agency one-pager
1	SMS	Demo nudge
3	Email	"How a 12-client agency cut delivery time 70%"
5	LinkedIn	Light touch
7	Email	Slot ask
10	SMS	Last SMS
14	Email	Pilot proposal
21	Email	90-day re-engage

6. Email templates

T+0 — Recap

```
Subject: AI workforce recap for {{agency}}

{{first_name}} -

Three things from today:

1) Mission Control Supervisor orchestrating across 3 test
   client workspaces.
2) Demo share preset minting a watermarked link for
   {{test_prospect}} in 10 seconds.
3) Runtime supervision panel showing Hermes / OpenClaw /
   Claude Code health per client.

Signed demo link (7-day):
{{signed_demo_link}}

14-day pilot, $1, on one of your test clients:
{{calendar_link}}

- {{rep_name}}
```

T+3 — Reference

```
Subject: {{reference_agency}} cut delivery time 70%

{{first_name}} -

{{reference_agency}}, 12 active clients, in {{city}}.

After 90 days on Workspace tier:
- New client onboarding: 4 weeks → 1 day
- Per-client cost dropped 65%
- Net new clients added: 6 in 90 days (was 2/qtr before)

Same agency shape as yours. Walk their workspace?
{{calendar_link}}

- {{rep_name}}
```

T+7

Subject: 15 minutes this week, {{first_name}}?

You said onboarding takes ~{{onboarding_weeks}} weeks today. At \${{client_arr}}/yr per client, every week of delay is real money.

\$1 pilot, 14 days. Pick one test client.

- {{slot_1}}
- {{slot_2}}
- {{slot_3}}

- {{rep_name}}

T+14

Subject: 14-day pilot for {{agency}} – proposal

{{first_name}} –

Pilot one-pager attached.

- Day 0: Kickoff + test client chosen
- Day 4: Test client workspace live, runtimes connected
- Day 5: First mission complete on test client
- Day 13: Delivery-time delta reported

\$1 today. \$799/mo on conversion (Workspace).

Sign: {{pilot_signup_link}}

- {{rep_name}}

7. SMS templates

T+1

Hi {{first_name}} – recap on the multi-client AI workforce for {{agency}}. Signed demo: {{short_link}}. STOP.

T+4

Following up on Baseline for {{agency}}. 15-min call this week? {{calendar_short}}. STOP.

T+8

Last nudge. Pilot proposal or shelve until next quarter? Reply BACK or STOP.

Reference ROI math

Active clients	Onboarding savings / yr	Per-client cost drop
5	15 onboarding weeks saved	-50%
10	30 weeks	-60%
25	75 weeks	-75%

Onboarding-week recovery typically converts to 4–8 net new clients per year per 25-client capacity, depending on funnel strength.